

# Distributed Support Vector Machines via the Alternating Direction Method of Multipliers

Anders Poirel,  
AM 229, Fall 2020

March 31, 2021

## 1 Alternating Descent Method of Multipliers

### 1.1 Motivation

The *Alternating Descent Method of Multipliers* (ADMM) algorithm is motivated by a desire to combine the best properties of two simpler algorithms while overcoming their respective limitations [1].

On one hand, the *dual descent* algorithm has subproblems that can be easily distributed. However, this comes at the cost of strong assumptions on the form of the problem. On the other hand, the *method of multipliers* (MM) modifies dual ascent to relax the required assumptions, but this sacrifices the ability to distribute parts of the algorithm.

### 1.2 Algorithm

The ADMM algorithm solves problems of the form

$$\begin{aligned} \min_{x,z} \quad & f(x) + g(z) \\ \text{subj. to} \quad & Ax + Bz = c \end{aligned} \tag{1}$$

where  $x \in \mathbb{R}^m, z \in \mathbb{R}^n, A \in \mathbb{R}^{m \times p}, B \in \mathbb{R}^{n \times p}, c \in \mathbb{R}^p$

Like MM, the ADMM algorithm in practice solves an equivalent augmented problem

$$\begin{aligned} \min_{x,z} \quad & f(x) + g(z) + \frac{\rho}{2} \|Ax + Bz - c\|_2^2 \\ \text{subj. to} \quad & Ax + Bz = c \end{aligned} \tag{2}$$

Indeed, any satisfiable pair  $x, z$  zeroes out the extra term in the objective, recovering problem (1).

The associated augmented Lagrangian is

$$L_\rho(x, z, \nu) = f(x) + g(z) + \nu^\top (Ax + Bz + c) + \frac{\rho}{2} \|Ax + Bz - c\|_2^2$$

The algorithm is given by

---

**Algorithm 1:** ADMM

---

```

for  $k = 1, 2, \dots$  do
   $x_{k+1} \leftarrow \operatorname{argmin}_x L_\rho(x, z_k, \nu_k)$ 
   $x_{k+1} \leftarrow \operatorname{argmin}_z L_\rho(x_{k+1}, z, \nu_k)$ 
   $\nu_{k+1} \leftarrow \nu_k + \rho(Ax_{k+1} + Bz_{k+1} - c)$ 
end

```

---

### 1.3 Properties of ADMM

#### Convergence

A major advantage of ADMM is that it converges in many common scenario where dual ascent does not. For instance, dual ascent does not converge if the objective is linear [1]. In fact if the problem satisfies the following assumptions

- (1)  $\operatorname{epi}(f)$  and  $\operatorname{epi}(g)$  are closed, nonempty convex sets
- (2) The Lagrangian of problem formulation (1) has a saddle point

Then that is, the points approach feasibility and the objective approaches its optimum

$$\begin{aligned} Ax_k + Bz_k &\rightarrow 0 \text{ as } k \rightarrow \infty \\ f(x_k) + g(z_k) &\rightarrow p^* \text{ as } k \rightarrow \infty \end{aligned}$$

These assumptions hold in many applications, including the example discussed in a later section [1].

#### Seperability

The function  $f$  is *seperable* if

$$f(x) = \sum_{i=1}^n f + i(x_i)$$

where  $x = (x_1, \dots, x_n)$  and  $x_i$  are subvectors of  $x$ . If either  $f$  or  $g$  in the ADMM problem (1) are seperable, then the  $x$  and  $z$ - minimization steps in the algorithm can be decomposed into independant problems [1], which allows the algorithm to conduct these updates in parallel, yielding distributed solvers for many problems.

## 2 Distributed ADMM

### Global variable consensus ADMM

The power of the ADMM becomes clear when the objective of the minimization problem is *additive*, i.e. when it is of the form

$$\min \sum_{i=1}^N f_i(x) \quad (3)$$

where each of the  $f_i$  are convex. In this situation, one can derive a version of the ADMM algorithm that can be distributed, which can yield a sizeable improvement in performance. To obtain a distributed algorithm, the objective function should be *seperable*

Now, to transform an additive objective into a seperable objective, the problem can be rewritten as a *global consensus problem* with “local variables” which are all constrained to be equal to the original variable  $x$  [2],

$$\begin{aligned} \min \quad & \sum_{i=1}^N f_i(x_i) \\ \text{subj. to} \quad & x_i = z \quad \forall i = 1, \dots, n \end{aligned} \quad (4)$$

In many situations, the objective has an extra term in  $x$ , which is rewritten in terms of  $z$

$$\begin{aligned} \min \quad & \sum_{i=1}^N f_i(x_i) + g(z) \\ \text{subj. to} \quad & x_i = z \quad \forall i = 1, \dots, n \end{aligned} \quad (5)$$

The augmented Lagrangian is then

$$L_\rho(x_1, \dots, x_N, z, \nu) = g(z) + \sum_{i=1}^N \left( f_i(x_i) + (\nu_i)^\top (x_i - z) + \frac{\rho}{2} \|x_i - z\|_2^2 \right)$$

which yields the following ADMM algorithm after simplifying the z-update step [1]:

---

**Algorithm 2:** Global variable consensus ADMM

---

```

for  $k = 1, 2, \dots$  do
     $(x_i)_{k+1} \leftarrow \operatorname{argmin}_{x_i} \left( f_i(x_i) + (\nu_i)_k^\top (x_i - z_k) + \frac{\rho}{2} \|x_i - z_k\|_2^2 \right)$ 
     $z_{k+1} \leftarrow \operatorname{argmin}_z \left( g(z) + \frac{N\rho}{2} \|z - \bar{x}_{k+1} - (1/\rho)\bar{\nu}_k\|_2^2 \right)$ 
     $(\nu_i)_{k+1} \leftarrow (\nu_i)_k + \rho((\beta_i)_k - z_{k+1})$ 
end

```

---

This algorithm performs  $N$  distinct  $x_i$  and  $\nu_i$ -updates at each step, each of which can be done in parallel as they are independent of one-another. The  $z$ -update is handled by a central process, and coordinates the solutions of the subproblems solved in the  $x, \nu$  updates.

## 2.1 Regularized Model Estimation

The ADMM form (1) is particularly natural for regularized model estimation problems, where the objective can be written as the sum of a “loss” and a “penalty” term, i.e.

$$\min_{\beta} \quad l(\beta, X, y) + r(\beta)$$

where  $X \in \mathbb{R}^{m \times p}$  is a matrix of training examples,  $y \in \mathbb{R}^m$  is the vector of training labels, and  $\beta$  is a vector of model parameters. If the loss function  $l$  and regularization function  $r$  are respectively convex, the substitution  $\beta = z$  immediately yields a problem in desired form (1)

$$\begin{aligned} \min_{\beta, z} \quad & l(\beta, X, y) + r(z) \\ \text{subj. to} \quad & \beta - z = 0 \end{aligned} \tag{6}$$

When  $l$  is separable, this problem can be solved using the global consensus ADMM algorithm by splitting the problem over blocks of data [1], such that

$$X = \begin{bmatrix} X_1 \\ \vdots \\ X_N \end{bmatrix}, \quad y = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix}$$

where  $X_i \in \mathbb{R}^{m_i \times p}$ ,  $y \in \mathbb{R}^m$ ,  $\sum_{i=1}^n m_i = m$ , and  $l_i$  is the loss functions on the  $i^{th}$  block of data. The algorithm is then, following the result for problems

of the form (5) and using uses the more compact scaled form described in [1],

---

**Algorithm 3:** Consensus global variable ADMM for regularized model estimation

---

```

for  $k = 1, 2, \dots$  do
     $(\beta_i)_{k+1} \leftarrow \operatorname{argmin}_{\beta_i} ((l_i(\beta_i, X_i, y_i) + (\rho/2)\|\beta_i - z_k + (u_i)_k\|_2^2)$ 
     $z_{k+1} \leftarrow \operatorname{argmin}_z \left( r(z) + \frac{N\rho}{2}\|z\|_2^2 - \bar{\beta}_{k+1} - \bar{u}_k \right)$ 
     $(u_i)_{k+1} \leftarrow (u_i)_k + (\beta_i)_k - z_{k+1}$ 
end

```

---

The next section demonstrates a worked out example of a straightforward application of this algorithm.

### 3 Support Vector Classifiers

#### 3.1 Problem formulation

Letting  $x_1, \dots, x_m \in \mathbb{R}^p$  be the training examples, and  $y_1, \dots, y_N \in \{-1, 1\}$  be the training labels, the standard form of the binary *support vector classifier* (SVC) problem is given in [3] by

$$\begin{aligned}
 & \operatorname{argmin}_{\beta, \beta_0, \xi} \quad \|\beta\|_2^2 \\
 & \text{subj. to} \quad y_i(x_i^\top \beta + \beta_0) \geq 1 - \xi \quad \forall i = 1, \dots, m \\
 & \quad \quad \quad \xi_i \geq 0 \quad \quad \quad \forall i = 1, \dots, m \\
 & \quad \quad \quad \sum_{i=1}^m \xi_i \leq C \quad \quad \quad C \text{ constant}
 \end{aligned} \tag{7}$$

Equivalently, we can formulate this as

$$\begin{aligned}
 & \operatorname{argmin}_{\beta, \beta_0, \xi} \quad \frac{1}{2}\|\beta\|_2^2 + C \sum_{i=1}^m \xi_i \\
 & \text{subj. to} \quad y_i(x_i^\top \beta + \beta_0) \geq 1 - \xi_i \quad \forall i = 1, \dots, m \\
 & \quad \quad \quad \xi_i \geq 0 \quad \quad \quad \forall i = 1, \dots, m
 \end{aligned} \tag{8}$$

Indeed, this minimum is defined iff  $\sum \xi_i$  is finite.

### 3.2 Motivation

In many potential applications of the SVC, there are a large number of training samples (say, millions), each of relatively modest dimensionality - i.e.  $N$  is much larger than  $p$ . While standard solvers for linear support vector classifiers perform well in this situation [4], the standard `libSVM` solvers for nonlinear kernels struggle on problems of this scale when the data is non-sparse [5].

The approach presented here uses the alternating descent method of multipliers to derive a parallelizable algorithm for fitting support vector classifiers, which will be capable of handling this type of large-scale problem. While only the linear classifier is shown here due to space constraints, extending this technique to non-linear kernels is straightforward [6].

### 3.3 Derivation of a distributed algorithm for SVC

The first goal is to rewrite the SVC problem (8) in a more convenient form to apply the global consensus variable ADMM algorithm. Formulation (8) is equivalent to

$$\begin{aligned} \operatorname{argmin}_{\beta, \beta_0, \xi} \quad & \frac{1}{2} \|\beta\|_2^2 + C \sum_{i=1}^m \xi_i \\ \text{subj. to} \quad & \xi_i = [1 - y_i(\beta^\top x_i + \beta_0)]_+ \end{aligned}$$

where  $[1 - y_i(\beta^\top x_i + \beta_0)]_+ = \max(0, 1 - y_i(\beta^\top x_i + \beta_0))$ . This is then equivalent to

$$\begin{aligned} \operatorname{argmin}_{\beta, \beta_0, \xi} \quad & \frac{1}{2C} \|\beta\|_2^2 + \sum_{i=1}^m \xi_i \\ \text{subj. to} \quad & \xi_i = [1 - y_i(\beta^\top x_i + \beta_0)]_+ \end{aligned}$$

Now, minimizing over  $\xi$ , this problem is equivalent to

$$\operatorname{argmin}_{\beta, \beta_0} \quad \sum_{i=1}^m [1 - y_i(\beta^\top x_i + \beta_0)]_+ + \frac{1}{2C} \|\beta\|_2^2 \quad (9)$$

Indeed, to minimize  $\sum \xi_i$  it suffices to take the smallest  $\xi_i$  allowed by the constraints, that is  $[1 - y_i(\beta^\top x_i + \beta_0)]_+$ . This corresponds to the *penalization*

method formulation of the SVC given in [3].

The problem is now in the desired “loss + penalty” form. The term  $\sum_{i=1}^n [1 - y_i(\beta^\top x_i + \beta_0)]_+$  is separable in  $x_1, \dots, x_m$ . Furthermore, both terms are convex in  $\beta$  as  $[1 - y_i(\beta^\top x_i + \beta_0)]_+$  is the pointwise maximum of two convex functions.

Thus substituting  $z = \beta$ ,

$$\begin{aligned} & \underset{\beta, \beta_0}{\operatorname{argmin}} \quad \sum_{i=1}^n [1 - y_i(\beta^\top x_i + \beta_0)]_+ + \frac{1}{2C} \|\beta\|_2^2 \\ & \text{subj. to} \quad x_i - z_i = 0 \quad \forall i = 1, \dots, m \end{aligned} \quad (10)$$

This problem can be solved using the global variable consensus ADMM algorithm, where  $X_i = [x_{i1} \ \dots \ x_{im_i}]^\top \in \mathbb{R}^{m_i \times p}$

---

**Algorithm 4:** Naive global variable consensus ADMM for SVC

---

```

for  $k = 1, 2, \dots$  do
   $(\beta_i)_{k+1} \leftarrow$ 
     $\underset{\beta_i}{\operatorname{argmin}} \left( \sum_{j=1}^{m_i} [1 - y_{ij}(\beta_i^\top x_{ij} + \beta_0)]_+ + (\rho/2) \|\beta_i - z_k + (u_i)_k\|_2^2 \right)$ 
   $z_{k+1} \leftarrow \underset{z}{\operatorname{argmin}} \left( \frac{1}{2C} \|z\|_2^2 + \frac{N\rho}{2} \|z\|_2^2 - \bar{\beta}_{k+1} - \bar{u}_k \right)$ 
   $(u_i)_{k+1} \leftarrow (u_i)_k + (\beta_i)_k - z_{k+1}$ 
end
```

---

This can be simplified further by solving the  $z$ -update analytically. Rewrite the right-hand side as

$$h(z) = \frac{1}{2C} z^\top I z + (z - \bar{\beta}_{k+1} - \bar{u}_k)^\top I (z - \bar{\beta}_{k+1} - \bar{u}_k)$$

Then, solve

$$\begin{aligned} 0 &= \nabla_z h(z) \\ 0 &= z^\top \left( \frac{1}{C} + N\rho \right) + N\rho (-\bar{\beta}_{k+1} - \bar{u}_k) \\ z &= \frac{N\rho}{1/C + N\rho} (\bar{\beta}_{k+1} + \bar{u}_k) \end{aligned}$$

This yields the final form of the ADMM algorithm for this problem,

---

**Algorithm 5:** Global variable consensus ADMM for SVC

---

```

for  $k = 1, 2, \dots$  do
     $(\beta_i)_{k+1} \leftarrow$ 
         $\operatorname{argmin}_{\beta_i} \left( \sum_{j=1}^{m_i} [1 - y_{ij}(\beta_i^\top x_{ij} + \beta_0)]_+ + (\rho/2) \|\beta_i - z_k + (u_i)_k\|_2^2 \right)$ 
     $z_{k+1} \leftarrow \frac{N\rho}{1/C + N\rho} (\bar{\beta}_{k+1} + \bar{u}_k)$ 
     $(u_i)_{k+1} \leftarrow (u_i)_k + (\beta_i)_k - z_{k+1}$ 
end

```

---

As in the general case, the  $u_i$ - and  $\beta_i$ -updates are parallelizable. Notice that the minimization problem in the  $\beta_i$  subproblems step resembles formulation (9) of the SVC. Indeed, one can treat this as a modified SVM problem and use existing single-process solvers [1].

This pattern of subproblems having the same form as the original problem appears frequently in applications of ADMM. In this sense, ADMM can be seen as a technique to extend nonparallel methods to large-scale problems [1].

## References

- [1] S. Boyd, N. Parikh, E. Chu, B. Peleato, and J. Eckstein, “Distributed optimization and statistical learning via the alternating direction method of multipliers,” *Found. Trends Mach. Learn.*, vol. 3, p. 1–122, Jan. 2011.
- [2] A. Nedic and A. Ozdaglar, “Cooperative distributed multi-agent optimization,” *Convex optimization in signal processing and communications*, vol. 340, 2010.
- [3] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer Science & Business Media, 2009.
- [4] R.-E. Fan, K.-W. Chang, C.-J. Hsieh, X.-R. Wang, and C.-J. Lin, “Liblinear: A library for large linear classification,” *Journal of machine learning research*, vol. 9, no. Aug, pp. 1871–1874, 2008.



- [5] C.-C. Chang and C.-J. Lin, “Libsvm: A library for support vector machines,” *ACM transactions on intelligent systems and technology (TIST)*, vol. 2, no. 3, pp. 1–27, 2011.
- [6] P. A. Forero, A. Cano, and G. B. Giannakis, “Consensus-based distributed support vector machines,” *Journal of Machine Learning Research*, vol. 11, no. 5, 2010.